

Introduction to pandas

https://bioinfmsc8.bio.ed.ac.uk/AY25_BPSM17.html

Quick links:
Click on any  to get back here!

- [venvs](#)
- [Series and dataframes](#)
- [Reading from a file](#)
- [DOT TAB TAB](#)
- [Slices](#)
- [head\(\)](#) and [tail\(\)](#)
- [describe\(\)](#)
- [NA values](#)
- [Selecting from dataframes](#)
- [Count unique values](#)
- [Filtering dataframes](#)
- [Subset the data](#)
- [Save as csv](#)
- [Comparison operators](#)
- [Multiple conditions](#)
- [apply\(\)](#) and [lambda](#)
- [Sorting dataframes](#), also on [multiple columns](#)

- [Creating and updating columns](#)
- [Descriptive statistics](#)
- [Indexing](#)
- [Selections using \[iloc\]\(#\) or the \[loc\]\(#\) indexer](#)
- [Some other cool things: \[agg\]\(#\) \[regate\]\(#\), \[groupby\]\(#\)](#)

Exercises

1. [how many fungal species have genomes bigger than 100Mb? What are their names?](#)
 2. [how many of each Kingdom/group \(plants, animals, fungi and protists\) have been sequenced?](#)
 3. [which *Heliconius* species genomes have been sequenced?](#)
 4. [which sequencing centre has sequenced the most plant genomes? the most insect genomes?](#)
 5. [add a column giving the number of proteins per gene. Which genomes have at least 10% more proteins than genes?](#)
-
- [important git/GitHub stuff at the end](#)

As always, there are many ways to do things in Python (and other languages), but [pandas](#) is almost always our first choice for dealing with tabular data objects in Python. If we have to process any data that are like Excel tables, pandas is probably [the](#) way.

It has its own web site: <https://pandas.pydata.org/>

The name [pandas](#) has little to do with the cute black and white bears! Two years ago, we could go see Yang Guang at the [Edinburgh Zoo](#), either by visiting the Zoo or watching the [live panda webcam](#)!


Last year, the attraction was [Haggis, the pygmy hippo baby/calf](#), now moved to [a different place!](#)

This year, starting yesterday, the attraction is a [female cheetah](#) that they hope will breed with the male they already have...

Meanwhile, back to Python, I believe the name is Pandas is derived from the term "Panel data", but ultimately, it really **doesn't** matter where the name came from...!

Today we are going to see how to:

- input some data into a dataframe
- manipulate it in various ways
- output subsets of the data



Getting started!

Like most Python modules or libraries, there's a common convention for importing:

Start python3

python3

```
Python 3.12.3 (main, Aug 14 2025, 17:47:21) [GCC 13.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
```

Import the modules we might need

```
import os, sys, subprocess
import numpy as np
import pandas as pd
```

```
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ModuleNotFoundError: No module named 'pandas'
```

If you find a module is missing, one option used to be to install a personal version!

subprocess.call("pip3 install pandas", shell = True)

```
error: externally-managed-environment
```

```
× This environment is externally managed
```

```
  ↳ To install Python packages system-wide, try apt install
    python3-xyz, where xyz is the package you are trying to
    install.
```

```
If you wish to install a non-Debian-packaged Python package,
create a virtual environment using python3 -m venv path/to/venv.
Then use path/to/venv/bin/python and path/to/venv/bin/pip. Make
sure you have python3-full installed.
```

```
If you wish to install a non-Debian packaged Python application,
it may be easiest to use pipx install xyz, which will manage a
virtual environment for you. Make sure you have pipx installed.
```

```
See /usr/share/doc/python3.12/README.venv for more information.
```

```
note: If you believe this is a mistake, please contact your Python installation or OS dist
hint: See PEP 668 for the detailed specification.
```

exit()

Try using **apt** to install it?

apt install pandas

```
E: Could not open lock file /var/lib/dpkg/lock-frontent - open (13: Permission denied)
E: Unable to acquire the dpkg frontend lock (/var/lib/dpkg/lock-frontent), are you root?
```

When **root** isnt around to help.....

[Help link 1](#)

[Help link 2](#)

How about using **pipx**?

pipx

```
usage: pipx [-h] [--quiet] [--verbose] [--version]
           {install,uninject,inject,upgrade,upgrade-all,uninstall,uninstall-all,reinstall
           ...
```

Install and execute **apps** from Python packages.

Binaries can either be installed globally into isolated Virtual Environments or run directly in a temporary Virtual Environment.

optional environment variables:

PIPX_HOME	Overrides default pipx location. Virtual Environments will be inst
PIPX_BIN_DIR	Overrides location of app installations. Apps are symlinked or cop
PIPX_MAN_DIR	Overrides location of manual pages installations. Manual pages are
PIPX_DEFAULT_PYTHON	Overrides default python used for commands.
USE_EMOJI	Overrides emoji behavior. Default value varies based on platform.

options:

-h, --help	show this help message and exit
--quiet, -q	Give less output. May be used multiple times corresponding to the
--verbose, -v	Give more output.
--version	Print version and exit

subcommands:

Get help for commands with pipx COMMAND --help

{install,uninject,inject,upgrade,upgrade-all,uninstall,uninstall-all,reinstall,reinstall	
install	Install a package
uninject	Uninstall injected packages from an existing Virtual Environment
inject	Install packages into an existing Virtual Environment
upgrade	Upgrade a package
upgrade-all	Upgrade all packages. Runs `pip install -U <pkgname>` for each pac
uninstall	Uninstall a package
uninstall-all	Uninstall all packages
reinstall	Reinstall a package
reinstall-all	Reinstall all packages
list	List installed packages
run	Download the latest version of a package to a temporary virtual en `__pypackages__` directory (experimental).
runpip	Run pip in an existing pipx-managed Virtual Environment
ensurepath	Ensure directories necessary for pipx operation are in your PATH
environment	Print a list of environment variables and paths used by pipx.
completions	Print instructions on enabling shell completions for pipx

Check things are set up correctly?

pipx ensurepath

Success! Added /localdisk/home/aivens2/.local/bin to the PATH environment variable.

Consider adding shell completions for pipx. Run 'pipx completions' for instructions.

You will need to open a new terminal or re-login for the PATH changes to take effect.

Otherwise pipx is ready to go! 🌟 🌟 🌟

Try installing something

pipx install ruff

installed package ruff 0.14.4, installed using Python 3.12.3

These apps are now globally available

- ruff

⚠️ Note: '/localdisk/home/aivens2/.local/bin' is not on your PATH environment variable. These apps will not be globally accessible until your PATH is updated. Run 'pipx ensurepath' to automatically add it, or manually modify your PATH in your sh config file (i.e. ~/.bashrc).

done! 🌟 🌟 🌟

Exit the terminal and then come back in

exit

How else might we have done this, given that it is `~/.bashrc` that is being modified.....?!

Quick check....

pipx ensurepath

/localdisk/home/aivens2/.local/bin is already in PATH.

⚠️ All pipx binary directories have been added to PATH. If you are sure you want to proceed

Otherwise pipx is ready to go! 🌟 🌟 🌟

Try our installed `ruff` programme!

ruff

Ruff: An extremely fast Python `linter` and code formatter.

Usage: ruff [OPTIONS] <COMMAND>

Commands:

check	Run Ruff on the given files or directories
rule	Explain a rule (or all rules)
config	List or describe the available configuration options
linter	List all supported upstream linters
clean	Clear any caches in the current directory and any subdirectories
format	Run the Ruff formatter on the given files or directories
server	Run the language server
analyze	Run analysis over Python source code
version	Display Ruff's version
help	Print this message or the help of the given subcommand(s)

```
Options:
  -h, --help      Print help
  -V, --version   Print version

Log levels:
  -v, --verbose   Enable verbose logging
  -q, --quiet     Print diagnostics, but nothing else
  -s, --silent    Disable all logging (but still exit with status code "1" upon detecting d

Global options:
  --config <CONFIG_OPTION> Either a path to a TOML configuration file (`pyproject.toml`
                           find in a `ruff.toml` configuration file) overriding a spe
                           option always take precedence over all configuration files
                           `--config`
  --isolated       Ignore all configuration files

For help with a specific command, see: `ruff help <command>`.
```

Let's try installing **pandas** now...

pipx install pandas

Note: Dependent package 'numpy' contains 2 apps

- f2py
- numpy-config

No apps associated with package pandas. Try again with '--include-deps' to include apps of dependent packages, which are listed above. If you are attempting to install a library, pipx should not be used. Consider using pip or a similar tool instead.

Ignored/missed their warning, tried with dependencies....

pipx install --include-deps pandas

⚠ File exists at /localdisk/home/aivens2/.local/bin/f2py and points to /localdisk/home/aivens2/.local/share/pipx/venvs/pandas/bin/f2py. Not modifying.
installed package pandas 2.3.3, installed using Python 3.12.3
These apps are now globally available

- numpy-config
- f2py (symlink missing or pointing to unexpected location)

done! 🌟 🌟 🌟

Are we all OK now?

p3

Python 3.12.3 (main, Aug 14 2025, 17:47:21) [GCC 13.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.

import pandas

Traceback (most recent call last):
File "<stdin>", line 1, in <module>
ModuleNotFoundError: No module named 'pandas'

...meh!... 😞

📌Need/use a virtual environment?

If we need a special set of packages, and/or they can't be globally installed for some reason, then we can use what is called a *virtual environment*, a "bubble" set up the way we want it all to be...

Let's be tidy!

```
mkdir ~/Exercises/Lecture17
```

```
cd ~/Exercises/Lecture17
```

Create a virtual environment called `playtime`

```
python -m venv playtime
```

What happened?!

```
ls -al playtime
```

```
total 4
drwxr-xr-x 5 aivens2 g_aivens2  98 Nov 10 09:27 .
drwxr-xr-x 3 aivens2 g_aivens2  30 Nov 10 09:27 ..
drwxr-xr-x 2 aivens2 g_aivens2 212 Nov 10 09:27 bin
drwxr-xr-x 3 aivens2 g_aivens2  32 Nov 10 09:27 include
drwxr-xr-x 3 aivens2 g_aivens2  32 Nov 10 09:27 lib
lrwxrwxrwx 1 aivens2 g_aivens2   3 Nov 10 09:27 lib64 -> lib
-rw-r--r-- 1 aivens2 g_aivens2 190 Nov 10 09:27 pyvenv.cfg
```

What does the folder structure look like?

```
tree | head -n20
```

```
.
├── playtime
│   ├── bin
│   │   ├── activate
│   │   ├── activate.csh
│   │   ├── activate.fish
│   │   ├── Activate.ps1
│   │   └── pip
│   ├── include
│   ├── lib
│   ├── lib64 -> lib
│   └── pyvenv.cfg
```

```
| | | | | pip3
| | | | | pip3.12
| | | | | python -> /usr/bin/python
| | | | | python3 -> python
| | | | | python3.12 -> python
| | | | |
| | | | | include
| | | | | | | python3.12
| | | | |
| | | | | lib
| | | | | | | python3.12
| | | | | | | | | site-packages
| | | | | | | | | | | pip
| | | | | | | | | | | | | __init__.py
```

How many files are there!?

ls -aIR | wc -l

1774

Does that not take up a whole pile of disk space?

du -hc --max-depth 1

```
15M    ./playtime
15M    .
15M    total
```

Can we successfully install pandas now!?

playtime/bin/pip install pandas

```
Collecting pandas
  Using cached pandas-2.3.3-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl.me
Collecting numpy>=1.26.0 (from pandas)
  Using cached numpy-2.3.4-cp312-cp312-manylinux_2_27_x86_64.manylinux_2_28_x86_64.whl.met
Collecting python-dateutil>=2.8.2 (from pandas)
  Using cached python_dateutil-2.9.0.post0-py2.py3-none-any.whl.metadata (8.4 kB)
Collecting pytz>=2020.1 (from pandas)
  Using cached pytz-2025.2-py2.py3-none-any.whl.metadata (22 kB)
Collecting tzdata>=2022.7 (from pandas)
  Using cached tzdata-2025.2-py2.py3-none-any.whl.metadata (1.4 kB)
Collecting six>=1.5 (from python-dateutil>=2.8.2->pandas)
  Using cached six-1.17.0-py2.py3-none-any.whl.metadata (1.7 kB)
Using cached pandas-2.3.3-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl (12.
Using cached numpy-2.3.4-cp312-cp312-manylinux_2_27_x86_64.manylinux_2_28_x86_64.whl (16.6
Using cached python_dateutil-2.9.0.post0-py2.py3-none-any.whl (229 kB)
Using cached pytz-2025.2-py2.py3-none-any.whl (509 kB)
Using cached tzdata-2025.2-py2.py3-none-any.whl (347 kB)
Using cached six-1.17.0-py2.py3-none-any.whl (11 kB)
Installing collected packages: pytz, tzdata, six, numpy, python-dateutil, pandas
Successfully installed numpy-2.3.4 pandas-2.3.3 python-dateutil-2.9.0.post0 pytz-2025.2 s
```

Try it now!

python

```
Python 3.12.3 (main, Aug 14 2025, 17:47:21) [GCC 13.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
```


import pandas

```
Traceback (most recent call last):  
  File "<stdin>", line 1, in <module>  
ModuleNotFoundError: No module named 'pandas'
```

D'oh!



DOH! Need to activate our **playtime** environment!

source playtime/bin/activate

```
(playtime) aivens2@bioinfmsc8:~/Exercises/Lecture17$
```

Try it now, this time we are in the environment!

python

```
Python 3.12.3 (main, Aug 14 2025, 17:47:21) [GCC 13.3.0] on linux  
Type "help", "copyright", "credits" or "license" for more information.
```

import pandas

pandas.__version__

'2.3.3'

Success!!

Exiting our environment once we are all done

(playtime) aivens2@bioinfmsc8:~/Exercises/Lecture17\$ deactivate

```
aivens2@bioinfmsc8:~/Exercises/Lecture17$
```

Try it again, this time we are NOT in the environment!

python

```
Python 3.12.3 (main, Aug 14 2025, 17:47:21) [GCC 13.3.0] on linux  
Type "help", "copyright", "credits" or "license" for more information.
```

import pandas

```
Traceback (most recent call last):  
  File "<stdin>", line 1, in <module>  
ModuleNotFoundError: No module named 'pandas'
```

Let's just get on with things....

source ~/Exercises/Lecture17/playtime/bin/activate

python

Python 3.12.3 (main, Aug 14 2025, 17:47:21) [GCC 13.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.

Import all the modules we want

```
import os, sys, subprocess
import numpy as np
import pandas as pd
```

📌 Series and dataframes

There are two types of things we need to know about: **Series** and **Dataframe** (name stolen from R). As with most things Python, there is lots of on-line documentation.

- <https://pandas.pydata.org/pandas-docs/stable/reference/series.html>
- <https://pandas.pydata.org/pandas-docs/stable/reference/frame.html>

A Series is like a list:

```
pd.Series(['a', 'b', 'c', 'd'])
```

```
0    a
1    b
2    c
3    d
dtype: object
```

However, unlike a list it has an explicit index (the thing on the left: 0, 1, 2, 3, essentially just the row numbers).

Also, the elements have a type (here it was **object** because they're strings).

If we put several Series objects together in a dict and assign them descriptive names then we can create a **DataFrame** object:

Make a dataframe

```
s1 = pd.Series(['a', 'b', 'c', 'd'])
s2 = pd.Series([2,4,8,16])
pd.DataFrame( { 'letter' : s1, 'number' : s2 } )
```

```
   letter  number
0      a        2
1      b        4
2      c        8
3      d       16
```

[Give feedback](#)

[Ask for help](#)

We still have an explicit index (0, 1, 2, 3), AND the dataframe displays nicely as a table.

📁 Creating dataframes from files

For data manipulation/visualisation, we generally create a dataframe by reading data from a file.

There are many ways to input/output data... described in great detail on <https://pandas.pydata.org/pandas-docs/stable/reference/io.html>

Today we are going to use an NCBI file that contains information about the various eukaryotic genome sequences they have.

We can get the file using `wget`:

```
# I did this on 10th Nov 2025, it's not changed since last year 🤖
```

```
subprocess.call('wget -qO eukaryotes.txt  
"ftp://ftp.ncbi.nlm.nih.gov/genomes/GENOME_REPORTS/eukaryotes.txt" ',  
shell=True)
```

```
0
```

```
os.system("wc -l eukaryotes.txt")
```

```
37951 eukaryotes.txt  
0
```

```
os.system("head -n3 eukaryotes.txt")
```

#Organism/Name	TaxID	BioProject	Accession	BioProject ID	Group	SubGroup	Size (Mb)	GC%	Assembly	Accession	Replicons	WGS	Scaffol
Emiliania huxleyi	CCMP1516	280463	PRJNA77753	77753	Protists	Other Protists	167.676	64.5	GCA_000372725.1	-	AHAL01	7795 38549	38554
Arabidopsis thaliana	3702	PRJNA10719	10719	Plants	Land Plants	119.669	36.0529	GCA_000001735.2	chromosome 1:NC_003070.9/CP002684.1; chromosome 2:NC_003071.1/CP002684.2				

First thing that we need to do is import the eukaryote genomes file into a dataframe table.

Notice that:

- the method we use is called `read_csv()` even though it's not actually a CSV file (it is actually tab-delimited, as we saw)
- we have to give the separator as a keyword argument
- pandas automatically treats the first line as a header row (column names)
- `df` is the conventional variable name for a dataframe
- we get an explicit index starting from zero

I need help with using read_csv

help(pd.read_csv)

Help on function read_csv in module pandas.io.parsers:

```
read_csv(filepath_or_buffer, sep=',', dialect=None, compression='infer', doublequote=True,
          Read CSV (comma-separated) file into DataFrame
```

Also supports optionally iterating or breaking of the file into chunks.

Additional help can be found in the `online docs for IO Tools <http://pandas.pydata.org/pandas-docs/stable/io.html>`.

Parameters

`filepath_or_buffer` : string or file handle / StringIO

The string could be a URL. Valid URL schemes include http, ftp, s3, and file. For file URLs, a host is expected. For instance, a local file could be file ://localhost/path/to/table.csv

`sep` : string, default ','

Delimiter to use. If sep is None, will try to automatically determine this. Regular expressions are accepted.

`engine` : {'c', 'python'}

Parser engine to use. The C engine is faster while the python engine is currently more feature-complete.

`lineterminator` : string (length 1), default None

Character to break file into lines. Only valid with C parser

`quotechar` : string (length 1)

The character used to denote the start and end of a quoted item. Quoted items can include the delimiter and it will be ignored.

`quoting` : int or csv.QUOTE_* instance, default None

Control field quoting behavior per ``csv.QUOTE\_\*`` constants. Use one of QUOTE_MINIMAL (0), QUOTE_ALL (1), QUOTE_NONNUMERIC (2) or QUOTE_NONE (3). Default (None) results in QUOTE_MINIMAL behavior.

`skipinitialspace` : boolean, default False

Skip spaces after delimiter

`escapechar` : string (length 1), default None

One-character string used to escape delimiter when quoting is QUOTE_NONE.

`dtype` : Type name or dict of column -> type, default None

... LOTS more ...

We can also use DOT TAB TAB method to what is available

pd.re # TAB TAB

pd.read_clipboard(	pd.read_feather(	pd.read_hdf(	pd.read_orc(	pd.read_sa
pd.read_csv(	pd.read_fwf(	pd.read_html(	pd.read_parquet(	pd.read_sp
pd.read_excel(	pd.read_gbq(	pd.read_json(	pd.read_pickle(	pd.read_sq


```
# Let's read the file in
```

```
df = pd.read_csv('eukaryotes.txt', sep="\t")
```

```
df
```

```
# Note the index on the left
```

```
      #Organism/Name  TaxID  ...  
0      Emiliana huxleyi CCMP1516  280463  ...  
1      Arabidopsis thaliana      3702  ...      The Arabidopsis Information Re  
2      Neopyropia yezoensis      2788  ...      Oce  
3      Glycine max      3847  ...      US DOE Joint Genome Instit  
4      Medicago truncatula      3880  ...  
...      ...      ...      ...  
37945      Saccharomyces cerevisiae NC-02  947038  ...      Genome Sequencing Center (GSC) at Wa  
37946      Saccharomyces cerevisiae CLIB382  947035  ...      Genome Sequencing Center (GSC) at Wa  
37947      Saccharomyces cerevisiae      4932  ...      Shanghai Jiao To  
37948      Saccharomyces cerevisiae      4932  ...      Chung-A  
37949      Saccharomyces cerevisiae      4932  ...      San  
  
[37950 rows x 19 columns]
```

```
# Similar things exist for "saving out" from a dataframe (once you have one!  
e.g. called df)
```

```
df.to_# TAB TAB
```

```
df.to_clipboard(  df.to_feather(  df.to_json(  df.to_parquet(  df.to_sql(  
df.to_csv(  df.to_gbq(  df.to_latex(  df.to_period(  df.to_stata(  
df.to_dict(  df.to_hdf(  df.to_markdown(  df.to_pickle(  df.to_string(  
df.to_excel(  df.to_html(  df.to_numpy(  df.to_records(  df.to_timestamp(
```

```
# How big is the data frame?
```

```
df.shape
```

```
(37950, 19)
```

```
# What are the column headings? Are the spaces an issue?
```

```
df.columns
```

```
Index(['#Organism/Name', 'TaxID', 'BioProject Accession', 'BioProject ID',  
      'Group', 'SubGroup', 'Size (Mb)', 'GC%', 'Assembly Accession',  
      'Replicons', 'WGS', 'Scaffolds', 'Genes', 'Proteins', 'Release Date',  
      'Modify Date', 'Status', 'Center', 'BioSample Accession'],  
      dtype='object')
```


list(df.columns)

```
['#Organism/Name', 'TaxID', 'BioProject Accession', 'BioProject ID', 'Group', 'SubGroup',
```



Get the DOT TAB TAB help!

df.

Display all 219 possibilities? (y or n)

df.Center	df.combine()	df.floordiv()	df.mean()	df.replace()	df.to_dict()
df.Genes	df.combine_first()	df.from_dict()	df.median()	df.resample()	df.to_excel()
df.Group	df.compare()	df.from_records()	df.melt()	df.reset_index()	df.to_feather()
df.Proteins	df.convert_dtypes()	df.ge()	df.memory_usage()	df.rfloordiv()	df.to_gbq()
df.Replicons	df.copy()	df.get()	df.merge()	df.rmod()	df.to_hdf()
df.Scaffolds	df.corr()	df.groupby()	df.min()	df.rmul()	df.to_html()
df.Status	df.corrwith()	df.gt()	df.mod()	df.rolling()	df.to_json()
df.SubGroup	df.count()	df.head()	df.mode()	df.round()	df.to_latex()
df.T	df.cov()	df.hist()	df.mul()	df.rpow()	df.to_markdown()
df.TaxID	df.cummax()	df.iat	df.multiply()	df.rsub()	df.to_numpy()
df.WGS	df.cummin()	df.idmax()	df.ndim	df.rtruediv()	df.to_orc()
df.abs()	df.cumprod()	df.idmin()	df.ne()	df.sample()	df.to_parquet()
df.add()	df.cumsum()	df.iloc	df.nlargest()	df.select_dtypes()	df.to_period()
df.add_prefix()	df.describe()	df.index	df.notna()	df.sem()	df.to_pickle()
df.add_suffix()	df.diff()	df.infer_objects()	df.notnull()	df.set_axis()	df.to_records()
df.agg()	df.div()	df.info()	df.nsmallest()	df.set_flags()	df.to_sql()
df.aggregate()	df.divide()	df.insert()	df.nunique()	df.set_index()	df.to_stata()
df.align()	df.dot()	df.interpolate()	df.pad()	df.shape	df.to_string()
df.all()	df.drop()	df.isetitem()	df.pct_change()	df.shift()	df.to_timestamp()
df.any()	df.drop_duplicates()	df.isin()	df.pipe()	df.size	df.to_xarray()
df.apply()	df.droplevel()	df.isna()	df.pivot()	df.skew()	df.to_xml()
df.applymap()	df.dropna()	df.isnull()	df.pivot_table()	df.sort_index()	df.transform()
df.asfreq()	df.dtypes	df.items()	df.plot()	df.sort_values()	df.transpose()
df.asof()	df.duplicated()	df.iterrows()	df.pop()	df.sparse	df.truediv()
df.assign()	df.empty	df.itertuples()	df.pow()	df.squeeze()	df.truncate()
df.astype()	df.eq()	df.join()	df.prod()	df.stack()	df.tx_convert()
df.at	df.equals()	df.keys()	df.product()	df.std()	df.tx_localize()
df.at_time()	df.eval()	df.kurt()	df.quantile()	df.style	df.unstack()
df.attrs	df.ewm()	df.kurtosis()	df.query()	df.sub()	df.update()
df.axes	df.expanding()	df.last()	df.radd()	df.subtract()	df.value_counts()
df.backfill()	df.explode()	df.last_valid_index()	df.rank()	df.sum()	df.values
df.between_time()	df.fill()	df.le()	df.rdiv()	df.swapaxes()	df.var()
df.bfill()	df.fillna()	df.loc	df.reindex()	df.swaplevel()	df.where()
df.bool()	df.filter()	df.lt()	df.reindex_like()	df.tail()	df.xs()
df.boxplot()	df.first()	df.map()	df.rename()	df.take()	
df.clip()	df.first_valid_index()	df.mask()	df.rename_axis()	df.to_clipboard()	
df.columns	df.flags	df.max()	df.reorder_levels()	df.to_csv()	



To examine bits of the dataframe we can use slices

df[3:7]

	#Organism/Name	TaxID	...	
3	Glycine max	3847	...	US DOE Joint Genome Institute (J
4	Medicago truncatula	3880	...	
5	Solanum lycopersicum	4081	...	Solanaceae Genomics
6	Hordeum vulgare subsp. vulgare	112509	...	Leibniz Institute of Plant Genetics and Cr

[4 rows x 19 columns]



Or use built in head() and tail() methods

df.head()

	#Organism/Name	TaxID	...	Center	B
0	Emiliania huxleyi CCMP1516	280463	...	JGI	
1	Arabidopsis thaliana	3702	...	The Arabidopsis Information Resource (TAIR)	
2	Neopyropia yezeensis	2788	...	Ocean University	
3	Glycine max	3847	...	US DOE Joint Genome Institute (JGI-PGF)	
4	Medicago truncatula	3880	...	NCBI	

[5 rows x 19 columns]



To get more information on the columns we have ended up with, use

```
describe()
```

```
df.describe()
```

	TaxID	BioProject ID	Size (Mb)	Scaffolds
count	3.795000e+04	3.795000e+04	37950.000000	3.795000e+04
mean	5.613396e+05	6.506686e+05	581.685245	5.681285e+04
std	8.356700e+05	2.512643e+05	1429.903851	2.871059e+05
min	2.708000e+03	1.190000e+02	0.000520	1.000000e+00
25%	9.606000e+03	4.838862e+05	28.917325	1.460000e+02
50%	1.194430e+05	6.691910e+05	58.418100	9.940000e+02
75%	7.495890e+05	8.332210e+05	661.516750	8.963000e+03
max	3.135897e+06	1.153517e+06	91113.700000	1.167394e+07

By default, this just shows us the numerical columns.

Why no stats for GC content, number of genes, etc?

Because in the original file, unknown values are represented with a hyphen "-", and pandas doesn't know that that means "missing data".

↑ We can tell it by using the `na_values` argument and giving a list of strings that mean "missing data".

Re-read the file in, telling pandas about "missing data"

```
df = pd.read_csv('eukaryotes.txt', sep="\t", na_values=['-'])
```

```
df.describe()
```

	TaxID	BioProject ID	Size (Mb)	GC%	Scaffolds	Genes
count	3.795000e+04	3.795000e+04	37950.000000	37376.000000	3.795000e+04	7.431000e+03
mean	5.613396e+05	6.506686e+05	581.685245	42.458719	5.681285e+04	1.822345e+04
std	8.356700e+05	2.512643e+05	1429.903851	7.666657	2.871059e+05	8.392855e+04
min	2.708000e+03	1.190000e+02	0.000520	0.134504	1.000000e+00	1.000000e+00
25%	9.606000e+03	4.838862e+05	28.917325	37.400000	1.460000e+02	9.007500e+03
50%	1.194430e+05	6.691910e+05	58.418100	41.400000	9.940000e+02	1.242000e+04
75%	7.495890e+05	8.332210e+05	661.516750	48.200000	8.963000e+03	1.771950e+04
max	3.135897e+06	1.153517e+06	91113.700000	78.373100	1.167394e+07	4.736081e+06

Now we can see all sorts of interesting info for GC content, number of scaffolds, etc. Specifying missing data in this way is very useful - it means that pandas will automatically skip values when calculating things AND when plotting...

🏠 Selecting from dataframes

What can we do with a dataframe once we have it!?

We can extract individual columns

```
df['Organism/Name']
```

```
Traceback (most recent call last):
  File "/localdisk/home/aivens2/Exercises/Lecture17/playtime/lib/python3.12/site-packages
    return self._engine.get_loc(casted_key)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
File "pandas/_libs/index.pyx", line 167, in pandas._libs.index.IndexEngine.get_loc
File "pandas/_libs/index.pyx", line 196, in pandas._libs.index.IndexEngine.get_loc
File "pandas/_libs/hashtable_class_helper.pxi", line 7088, in pandas._libs.hashtable.PyO
File "pandas/_libs/hashtable_class_helper.pxi", line 7096, in pandas._libs.hashtable.PyO
KeyError: 'Organism/Name'
```

The above exception was the direct cause of the following exception:

```
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  File "/localdisk/home/aivens2/Exercises/Lecture17/playtime/lib/python3.12/site-packages/
    indexer = self.columns.get_loc(key)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/localdisk/home/aivens2/Exercises/Lecture17/playtime/lib/python3.12/site-packages/
    raise KeyError(key) from err
KeyError: 'Organism/Name'
```

Get the name right!

```
df['#Organism/Name']
```

```
0      Emiliana huxleyi CCMP1516
1      Arabidopsis thaliana
2      Neopyropia yezoensis
3      Glycine max
4      Medicago truncatula
      ...
37945    Saccharomyces cerevisiae NC-02
37946    Saccharomyces cerevisiae CLIB382
37947    Saccharomyces cerevisiae
37948    Saccharomyces cerevisiae
37949    Saccharomyces cerevisiae
Name: #Organism/Name, Length: 37950, dtype: object
```

Note the syntax: it looks like we're getting a single value from a dict called df, but

actually we're asking for that **column** from all rows. The result of doing this is a **Series** object.

📌 We can count unique (non-redundant) values in a series and pandas will helpfully put the most frequent at the top.

```
df['#Organism/Name'].value_counts()
Saccharomyces cerevisiae      1412
Homo sapiens                  1215
Fusarium oxysporum            398
Pyricularia oryzae            375
Aspergillus fumigatus         333
...
Tilapia sparmanii             1
Lamprologus tigris pictilis    1
Coptodon rendalli             1
Heterochromis multidentatus    1
Saccharomyces cerevisiae CLIB382 1
Name: #Organism/Name, Length: 17177, dtype: int64
```

How many were there?

```
len(df['#Organism/Name'].value_counts())
```

```
17177
```

If we give a **list** of column names instead of a single one then we get back a `dataframe` object.

For subsequent examples we will generally use `head()` to make it easier to see the output, but all the rows are really there!

```
df[ ['#Organism/Name', 'Size (Mb)', 'GC%'] ].head()
```

	#Organism/Name	Size (Mb)	GC%
0	Emiliana huxleyi CCMP1516	167.676	64.5000
1	Arabidopsis thaliana	119.669	36.0529
2	Neopyropia yezoensis	107.591	64.8454
3	Glycine max	978.942	35.1221
4	Medicago truncatula	430.008	33.4462

📌 Filtering dataframes

To figure out which rows satisfy some condition, we can use a special shorthand:

```
df['#Organism/Name'] == 'Arabidopsis thaliana'
```

```
0      False
1       True
2      False
3      False
4      False
...
37945  False
37946  False
37947  False
37948  False
37949  False
```

```
Name: #Organism/Name, Length: 37950, dtype: bool
```

So how many of each did we get?

```
(df['#Organism/Name'] == 'Arabidopsis thaliana').value_counts()
```

```
False    37707
True       243
```

```
Name: #Organism/Name, dtype: int64
```

Or use `isin`

```
df[df['#Organism/Name'].isin(["Arabidopsis thaliana"])]
```

```
.shape
(243, 19)
```

Be specific with your query!

```
df[df['#Organism/Name'].isin(["Arabidopsis"])]
```

```
.shape
(0, 19)
```

This looks like we are comparing the whole column to the string, but pandas uses some special "Python magic" to make it compare each row and return a Series of True/False values. This isn't very useful on its own, but becomes useful when we slice the dataframe using those values, as shown below...



Make a subset of the data

```
cress = df[df['#Organism/Name'] == 'Arabidopsis thaliana']
```

cress

	#Organism/Name	TaxID	...	Center
1	Arabidopsis thaliana	3702	...	The Arabidopsis Information Resource (TAIR)
20076	Arabidopsis thaliana	3702	...	max planck institute for biology tuebingen
20080	Arabidopsis thaliana	3702	...	Community-Consensus Arabidopsis Thaliana Refer...
23218	Arabidopsis thaliana	3702	...	The Salk Institute for Biological Studies
23224	Arabidopsis thaliana	3702	...	max planck institute for biology tuebingen
...	...	...	...	...
34682	Arabidopsis thaliana	3702	...	MPI FOR DEVELOPMENTAL BIOLOGY
34683	Arabidopsis thaliana	3702	...	MPI FOR DEVELOPMENTAL BIOLOGY
34698	Arabidopsis thaliana	3702	...	MPI FOR DEVELOPMENTAL BIOLOGY
34699	Arabidopsis thaliana	3702	...	MPI FOR DEVELOPMENTAL BIOLOGY
34711	Arabidopsis thaliana	3702	...	MPI FOR DEVELOPMENTAL BIOLOGY

[243 rows x 19 columns]

⬆️ Notice that the result is itself a dataframe, which we could save out as a csv file.
See this for more details: https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.to_csv.html#pandas.DataFrame.to_csv

```
cress.to_csv("All_the_Arabidopsis_ones.tsv",sep="\t",header=True)
```

```
os.system("head -1 *; wc -l *")
```

```
==> All_the_Arabidopsis_ones.tsv <==
#Organism/Name TaxID BioProject Accession BioProject ID Group SubGroup

==> eukaryotes.txt <==
#Organism/Name TaxID BioProject Accession BioProject ID Group SubGroup Si
244 All_the_Arabidopsis_ones.tsv
37951 eukaryotes.txt
38195 total
0
```

We can filter/extract to ask questions like:

What are the sizes of all the Arabidopsis thaliana genomes?

```
df[df['#Organism/Name'] == 'Arabidopsis thaliana']['Size (Mb)']
```

1	119.669000
20076	143.501000
20080	142.481000
23218	132.081000
23224	139.933000
...	...
34682	2.035040
34683	0.844674
34698	0.838703
34699	0.325150

```
34711      0.847842
Name: Size (Mb), Length: 243, dtype: float64
```

Store the sizes as a Series object, then describe() it

```
sizes = df[df['#Organism/Name'] == 'Arabidopsis thaliana']['Size (Mb)']
```

sizes

```
1      119.669000
20076   143.501000
20080   142.481000
23218   132.081000
23224   139.933000
...
34682     2.035040
34683     0.844674
34698     0.838703
34699     0.325150
34711     0.847842
Name: Size (Mb), Length: 243, dtype: float64
```

sizes.describe()

```
count    243.000000
mean      94.320730
std       60.632455
min        0.195283
25%        2.020030
50%       131.229000
75%       136.562000
max       244.583000
Name: Size (Mb), dtype: float64
```

Other comparison operators work too!

```
df[df['Size (Mb)'] > 10000]
```

```

#Organism/Name  TaxID  ...  Cent
9      Triticum aestivum  4565  ...  NC
587    Neoceratodus forsteri  7892  ...  University of Konsta
609    Lepidosiren paradoxa  7883  ...  The Max Planck Institute of Molecular Cell Bio.
667      Bombina bombina  8345  ...  NC
668    Ichthyophis bannanicus  8453  ...  Northwestern Polytechnical Universi
...
37511    Homo sapiens  9606  ...  UNIVERSITY OF LEE
37515    Homo sapiens  9606  ...  UNIVERSITY OF LEE
37520    Homo sapiens  9606  ...  UNIVERSITY OF LEE
37523    Homo sapiens  9606  ...  UNIVERSITY OF LEE
37528    Homo sapiens  9606  ...  UNIVERSITY OF LEE

[112 rows x 19 columns]
```



Use & to join multiple conditions

note the brackets around the conditions

```
df[ (df['Size (Mb)'] > 10000) & (df['Status'] == 'Scaffold') ]
```


	#Organism/Name	...	BioSample Accession
1182	Pinus taeda	...	SAMN02981512
1195	Picea abies	...	SAMEA3269885
1215	Picea glauca	...	SAMN02736787
1364	Pseudotsuga menziesii	...	SAMN03333061
1385	Pinus lambertiana	...	SAMN03354659
1530	Sequoia sempervirens	...	SAMN11654426
1603	Picea sitchensis	...	SAMN04296985
4429	Larix sibirica	...	SAMN07326255
5415	Picea engelmannii	...	SAMN10388286
6901	Picea mariana	...	SAMN14324163
11866	Iris pallida	...	SAMN28206045
14473	Calotriton arnoldi	...	SAMEA113426058
17354	Sambucus nigra	...	SAMEA7522178
18293	Viscum album	...	SAMEA7522516
19320	Picea glauca	...	SAMN01120252
20770	Triticum dicoccoides	...	SAMN13041385
22734	Triticum dicoccoides	...	SAMEA6062056
22872	Allium cepa var. aggregatum	...	SAMD00071601
25477	Bubalus bubalis bubalis	...	SAMEA8110765
28559	Triticum aestivum	...	SAMN41075130
28826	Triticum aestivum	...	SAMN41710599
29058	Triticum aestivum	...	SAMEA6062055
29065	Triticum aestivum	...	SAMEA104424263
29279	Triticum aestivum	...	SAMEA3663800
29287	Triticum aestivum	...	SAMEA6374024
29479	Triticum aestivum	...	SAMEA6374020
29490	Triticum aestivum	...	SAMEA6374022
29668	Triticum aestivum	...	SAMEA6374021
29678	Triticum aestivum	...	SAMEA6374023
30023	Triticum aestivum	...	SAMEA3269884
37447	Homo sapiens	...	SAMEA8949006
37452	Homo sapiens	...	SAMEA8948994
37455	Homo sapiens	...	SAMEA8949015
37463	Homo sapiens	...	SAMEA8949021
37464	Homo sapiens	...	SAMEA8949033
37468	Homo sapiens	...	SAMEA8949024
37471	Homo sapiens	...	SAMEA8949031
37475	Homo sapiens	...	SAMEA8949034
37480	Homo sapiens	...	SAMEA8949004
37488	Homo sapiens	...	SAMEA8948990
37511	Homo sapiens	...	SAMEA8949025
37515	Homo sapiens	...	SAMEA8949022
37520	Homo sapiens	...	SAMEA8948980
37523	Homo sapiens	...	SAMEA8948989
37528	Homo sapiens	...	SAMEA8949007

[45 rows x 19 columns]

🏠 Your new best friend: the `apply()` method!

To run arbitrary expressions on rows, use the `apply()` method with a `lambda`.

What is a lambda?! Good description and examples here:

<https://www.afternerd.com/blog/python-lambdas/>.

Easiest way to think of it is that it acts like a "mini-function" that we can use without having to write a bigger "proper" function.

`axis=1` means that we want to iterate over rows rather than columns.

The argument to the lambda expression will be a single row so we can treat it as a `dict` to get individual columns

```
# Make all the organism names upper case
```

```
df.apply(lambda x : x['#Organism/Name'].upper(), axis=1).head()
```

```
0    EMILIANIA HUXLEYI CCMP1516
1      ARABIDOPSIS THALIANA
2    NEOPYROPIA YEZOENSIS
3           GLYCINE MAX
4    MEDICAGO TRUNCATULA
dtype: object
```

This comes in very handy for filtering with complex conditions!

```
# Find birds and mammals with "small" genome assemblies
```

```
# Which column do I need to look in?
```

```
df.columns
```

```
Index(['#Organism/Name', 'TaxID', 'BioProject Accession', 'BioProject ID',
      'Group', 'SubGroup', 'Size (Mb)', 'GC%', 'Assembly Accession',
      'Replicons', 'WGS', 'Scaffolds', 'Genes', 'Proteins', 'Release Date',
      'Modify Date', 'Status', 'Center', 'BioSample Accession'],
      dtype='object')
```

```
'Birds' in set(df['Group'])
```

False

'Birds' in set(df['SubGroup'])

True

'Mammals' in set(df['SubGroup'])

True

df[df.apply(lambda x : x['Size (Mb)'] < 500 and x['SubGroup'] in ['Birds'], axis=1)]

	#Organism/Name	TaxID	...	Center	BioSample Accessio
13438	Pyrrhura rhodoccephala	311893	...	Iridian Genomes	SAMN2558240
15504	Luscinia luscinia	391696	...	Charles University	SAMN2609580
15505	Luscinia megarhynchos	383689	...	Charles University	SAMN2609580
17605	Strigops habroptila	2489341	...	Vertebrate Genomes Project	SAMN0994875
19254	Ficedula albicollis	59894	...	Uppsala University	SAMEA316639
20151	Vidua chalybeata	81927	...	Vertebrate Genomes Project	SAMN1225400
22467	Calidris pugnax	198806	...	Uppsala University	SAMN2668815

[7 rows x 19 columns]

df[df.apply(lambda x : x['Size (Mb)'] < 500 and x['SubGroup'] in ['Birds'], axis=1)].shape[0]

7

df[df.apply(lambda x : x['Size (Mb)'] < 500 and x['SubGroup'] in ['Birds'], axis=1)].shape[1]

19

df[df.apply(lambda x : x['Size (Mb)'] < 500 and x['SubGroup'] in ['Mammals'], axis=1)].shape[0]

514

birds_and_mammals = df[df.apply(lambda x : x['Size (Mb)'] < 500 and x['SubGroup'] in ['Birds', 'Mammals'], axis=1)]

birds_and_mammals

	#Organism/Name	TaxID	...	Center	BioSampl
564	Osphranter robustus	9319	...	AUSTRALIAN NATIONAL UNIVERSITY	S.
578	Isodon macrourus macrourus	120737	...	AUSTRALIAN NATIONAL UNIVERSITY	S.
670	Phascogale tapoatafa	9293	...	AUSTRALIAN NATIONAL UNIVERSITY	S.
671	Potorous tridactylus	9310	...	AUSTRALIAN NATIONAL UNIVERSITY	S.
672	Pseudocheirus peregrinus	9333	...	AUSTRALIAN NATIONAL UNIVERSITY	S.
...	...	...	...	...	...
37808	Homo sapiens	9606	...	STANFORD UNIVERSITY	S.

Give feedback
Ask for help

37811	Homo sapiens	9606	...	Harvard Medical School	S.
37812	Homo sapiens	9606	...	Baylor College of Medicine	S.
37814	Homo sapiens	9606	...	Pacific Biosciences	S.
37815	Homo sapiens	9606	...	Telomere-to-Telomere Consortium	S.
[521 rows x 19 columns]					

birds_and_mammals.shape
(521, 19)
birds_and_mammals.shape[0]
521

📈 Sorting dataframes

To sort, there's a `sort_values()` method which mostly does what we want!

Which are the largest genomes?

```
df.sort_values('Size (Mb)', ascending=False).head()
```

	#Organism/Name	TaxID	...		Cent
18293	Viscum album	3972	...	WELLCOME SANGER INSTITUT	
609	Lepidosiren paradoxa	7883	...	The Max Planck Institute of Molecular Cell Bio..	
19655	Lepidosiren paradoxa	7883	...	The Max Planck Institute of Molecular Cell Bio..	
1739	Protopterus annectens	7888	...		NCB
587	Neoceratodus forsteri	7892	...		University of Konstan

[5 rows x 19 columns]

Notice that this returns a new dataframe (that we could have kept as a separate dataframe by assigning it to a new variable). Unlike lists, it doesn't affect the original (unless we pass `inplace=True`).

```
df.sort_values('Size (Mb)', ascending=False, inplace=True)
```

df

	#Organism/Name	TaxID	...		Cent
18293	Viscum album	3972	...	WELLCOME SANGER INSTITUT	
609	Lepidosiren paradoxa	7883	...	The Max Planck Institute of Molecular Cell Bio..	
19655	Lepidosiren paradoxa	7883	...	The Max Planck Institute of Molecular Cell Bio..	
1739	Protopterus annectens	7888	...		NCB
587	Neoceratodus forsteri	7892	...		University of Konstan
...	...	...	...		..
25442	Ornithodoros moubata	6938	...	Friedrich-Loeffler-Institu	
24448	Ornithodoros moubata	6938	...	Friedrich-Loeffler-Institu	
23448	Ornithodoros moubata	6938	...	Friedrich-Loeffler-Institu	
24810	Ornithodoros moubata	6938	...	Friedrich-Loeffler-Institu	
26358	Ornithodoros moubata	6938	...	Friedrich-Loeffler-Institu	

[37950 rows x 19 columns]

```
df['Size (Mb)']
```

18293	91113.700000
609	87218.800000
19655	46159.900000

```
1739      40054.300000
587       34557.600000
...
25442      0.001600
24448      0.001196
23448      0.000548
24810      0.000545
26358      0.000520
Name: Size (Mb), Length: 37950, dtype: float64
```



We can sort on multiple columns if we want

sort by name, then within each name with highest GC first

df.sort_values(['#Organism/Name', 'GC%'], ascending=[True, False]).head()

```
      #Organism/Name  ...  BioSample Accession
6194  Aaosphaeria arxii CBS 175.79  ...  SAMN05661102
11711  Aaosphaeria pasadenensis  ...  SAMN25226780
1795   Abelmoschus esculentus  ...  SAMN35805024
4315   Abeoforma whisleri  ...  SAMN06200030
8285   Abisara bifasciata  ...  SAMN12776983
```

```
[5 rows x 19 columns]
```

If we need to do something more complicated than this then we probably need to add a new column and then sort on it....

🏠 Creating and updating columns

For simple cases, we can create a new column by assigning values to it using the syntax we've already seen:

add a column for AT content

```
df['at content'] = 1 - (df['GC%'] / 100)
```

```
# view the new column
```

```
df[ ['#Organism/Name', 'GC%', 'at_content'] ].head()
```

	#Organism/Name	GC%	at_content
18293	Viscum album	31.2000	0.688000
609	Lepidosiren paradoxa	40.4307	0.595693
19655	Lepidosiren paradoxa	40.3000	0.597000
1739	Protopterus annectens	40.4512	0.595488
587	Neoceratodus forsteri	44.9367	0.550633

This changes the dataframe in place.

Having got used to referring to columns like this, this should work?

make a genus column with the first part of the species name

```
df['genus'] = df['#Organism/Name'].split(' ')[0]
```

```
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  File "/localdisk/home/aivens2/Exercises/Lecture17/playtime/lib/python3.12/site-packages/
    return object.__getattr__(self, name)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AttributeError: 'Series' object has no attribute 'split'. Did you mean: 'plot'?
```

but it doesn't - the "Python magic" only works on certain operators (mostly mathematical stuff).

`apply()` can help us out here.

```
df['genus'] = df.apply(lambda x : x['#Organism/Name'].split(' ')[0], axis=1)
```

df['#Organism/Name', 'genus'].head()

Traceback (most recent call last):

```
File "/localdisk/home/aivens2/Exercises/Lecture17/playtime/lib/python3.12/site-packages/pandas/core/indexes/base.py", line 100, in get_loc
    return self._engine.get_loc(casted_key)
          ~~~~~^~~~~~
File "pandas/_libs/index.pyx", line 167, in pandas._libs.index.IndexEngine.get_loc
File "pandas/_libs/index.pyx", line 196, in pandas._libs.index.IndexEngine.get_loc
File "pandas/_libs/hashtable_class_helper.pxi", line 7088, in pandas._libs.hashtable.PyObjectHashTable.get_item
File "pandas/_libs/hashtable_class_helper.pxi", line 7096, in pandas._libs.hashtable.PyObjectHashTable.get_item
KeyError: ('#Organism/Name', 'genus')
```

The above exception was the direct cause of the following exception:

Traceback (most recent call last):

```
File "<stdin>", line 1, in <module>
File "/localdisk/home/aivens2/Exercises/Lecture17/playtime/lib/python3.12/site-packages/pandas/core/indexes/base.py", line 100, in get_loc
    indexer = self.columns.get_loc(key)
              ~~~~~^~~~~~
File "/localdisk/home/aivens2/Exercises/Lecture17/playtime/lib/python3.12/site-packages/pandas/core/indexes/base.py", line 100, in get_loc
    raise KeyError(key) from err
KeyError: ('#Organism/Name', 'genus')
```

If you want to view columns, they need to be formatted as a list!

df[['#Organism/Name', 'genus']].head()

	#Organism/Name	genus
18293	Viscum album	Viscum
609	Lepidosiren paradoxa	Lepidosiren
19655	Lepidosiren paradoxa	Lepidosiren
1739	Protopterus annectens	Protopterus
587	Neoceratodus forsteri	Neoceratodus

📈 Descriptive statistics

We can take various measurements of a Series

```
print(df[ ['Size (Mb)', 'GC%'] ].mean())
```

```
Size (Mb)    581.685245  
GC%         42.458719  
dtype: float64
```

```
print(df[ ['Size (Mb)', 'GC%'] ].median())
```

```
Size (Mb)    58.4181  
GC%         41.4000  
dtype: float64
```

```
print(df[ ['Size (Mb)', 'GC%'] ].std())
```

```
Size (Mb)    1429.903851  
GC%         7.666657  
dtype: float64
```

Pearson correlation between columns too!

```
df[ ['Size (Mb)', 'Genes', 'Proteins', 'GC%'] ].corr()
```

	Size (Mb)	Genes	Proteins	GC%
Size (Mb)	1.000000	0.121138	0.471012	-0.134400
Genes	0.121138	1.000000	0.206491	-0.052944
Proteins	0.471012	0.206491	1.000000	-0.209682
GC%	-0.134400	-0.052944	-0.209682	1.000000

📌Indexing

So far we have not really talked about indexing, but it turns out to be quite important. The index is the thing on the left of the output: currently just the row number assigned when we first read the dataframe in.

df.head()

	#Organism/Name	TaxID	BioProject	Accession	...	BioSample	Accession	at_cont
18293	Viscum album	3972		PRJEB64803	...	SAMEA7522516		0.688
609	Lepidosiren paradoxa	7883		PRJNA808322	...	SAMN26083907		0.595
19655	Lepidosiren paradoxa	7883		PRJNA808321	...	SAMN26083907		0.597
1739	Protopterus annectens	7888		PRJNA754246	...	SAMN17890435		0.595
587	Neoceratodus forsteri	7892		PRJNA644903	...	SAMN15471169		0.550

[5 rows x 21 columns]

If there's a unique identifier somewhere in the data, then we can use that as the index by passing `index_col` when we read the data in.

Let's read PART OF IT in again, but specify what we want as the index

```
df_mini = pd.read_csv('eukaryotes.txt', sep="\t", na_values=['-'],  
index_col='Assembly Accession').head()
```

df_mini

	#Organism/Name	...	BioSample	Accession
Assembly Accession		...		
GCA_000372725.1	Emiliana huxleyi CCMP1516	...	SAMN02744062	
GCA_000001735.2	Arabidopsis thaliana	...	SAMN03081427	
GCA_009829735.1	Neopyropia yezoensis	...	SAMN13316713	
GCA_000004515.5	Glycine max	...	SAMN00002965	
GCA_003473485.2	Medicago truncatula	...	SAMN08400029	

[5 rows x 18 columns]

We could also use a non-unique index without any issues

```
df['#Organism/Name'].value_counts().head()
```

```
Homo sapiens          1060
Saccharomyces cerevisiae 1412
Homo sapiens          1215
Fusarium oxysporum     398
Pyricularia oryzae     375
Aspergillus fumigatus  333
Name: #Organism/Name, dtype: int64
```

```
df_org = pd.read_csv('eukaryotes.txt', sep="\t", na_values=['-'],
index_col='#Organism/Name')
```

```
df_org.head(2)
```

```

#Organism/Name          TaxID  ... BioSample Accession
Emiliana huxleyi CCMP1516  280463  ...          SAMN02744062
Arabidopsis thaliana      3702  ...          SAMN03081427

[2 rows x 18 columns]
```

```
# Everything from the original file is still there...
```

```
df_org.shape
```

```
(37950, 18)
```

```
# Different one this time, indexing on sequencing Center
```

```
df_center = pd.read_csv('eukaryotes.txt', sep="\t", na_values=['-'],
index_col='Center')
```

```
df_center.head(16)
```

```

#Organism/Name  ...  B
Center
JGI              Emiliana huxleyi CCMP1516  ...
The Arabidopsis Information Resource (TAIR)  Arabidopsis thaliana  ...
Ocean University  Neopyropia yezoensis  ...
US DOE Joint Genome Institute (JGI-PGF)      Glycine max  ...
NCBI              Medicago truncatula  ...
Solanaceae Genomics Project      Solanum lycopersicum  ...
Leibniz Institute of Plant Genetics and Crop Pl...  Hordeum vulgare subsp. vulgare  ...
NCBI              Oryza sativa Japonica Group  ...
MaizeGDB              Zea mays  ...
NCBI              Triticum aestivum  ...
Broad Institute      Pneumocystis carinii B80  ...
S. pombe European Sequencing Consortium (EUPOM)  Schizosaccharomyces pombe  ...
Saccharomyces Genome Database      Saccharomyces cerevisiae S288C  ...
NCBI              Aspergillus nidulans FGSC A4  ...
J. Craig Venter Institute      Aspergillus fumigatus Af293  ...
Broad Institute      Neurospora crassa OR74A  ...

[16 rows x 18 columns]
```

```
# Is the column still there in the dataframe?
```

```
df_center[ ['Center'] ]
```

```
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  File "/localdisk/home/s0000000/.local/lib/python3.8/site-packages/pandas/core/frame.py",
    indexer = self.loc._get_listlike_indexer(key, axis=1)[1]
  File "/localdisk/home/s0000000/.local/lib/python3.8/site-packages/pandas/core/indexing.py",
    self._validate_read_indexer(keyarr, indexer, axis)
  File "/localdisk/home/s0000000/.local/lib/python3.8/site-packages/pandas/core/indexing.py",
    raise KeyError(f"None of [{key}] are in the [{axis name}]")
KeyError: "None of [Index(['Center'], dtype='object')] are in the [columns]"
```

That was why there were 18 columns, instead of 19!

df_center.columns

```
Index(['#Organism/Name', 'TaxID', 'BioProject Accession', 'BioProject ID',
      'Group', 'SubGroup', 'Size (Mb)', 'GC%', 'Assembly Accession',
      'Replicons', 'WGS', 'Scaffolds', 'Genes', 'Proteins', 'Release Date',
      'Modify Date', 'Status', 'BioSample Accession'],
      dtype='object')
```

Rather than reading in the data again, we could set the index to be one of the existing columns...

df_rejigged = df_center.set_index('BioSample Accession')

df_rejigged

	#Organism/Name	TaxID	...	Modify Date	Status
BioSample Accession			...		
SAMN02744062	Emiliana huxleyi CCMP1516	280463	...	2014/08/01	Scaffold
SAMN03081427	Arabidopsis thaliana	3702	...	2022/10/20	Chromosome
SAMN13316713	Neopyropia yezoensis	2788	...	2020/01/06	Chromosome
SAMN00002965	Glycine max	3847	...	2021/03/10	Chromosome
SAMN08400029	Medicago truncatula	3880	...	2018/09/12	Chromosome
...	...	...	...	...	...
SAMN14365262	Saccharomyces cerevisiae NC-02	947038	...	2020/06/23	Scaffold
SAMN00198998	Saccharomyces cerevisiae CLIB382	947035	...	2014/08/04	Scaffold
SAMN20308677	Saccharomyces cerevisiae	4932	...	2023/03/27	Chromosome
SAMN26740367	Saccharomyces cerevisiae	4932	...	2024/05/24	Chromosome
SAMEA3138293	Saccharomyces cerevisiae	4932	...	2009/07/02	Chromosome

[37950 rows x 17 columns]

... however ...

df_rejigged.columns

```
Index(['#Organism/Name', 'TaxID', 'BioProject Accession', 'BioProject ID',
      'Group', 'SubGroup', 'Size (Mb)', 'GC%', 'Assembly Accession',
      'Replicons', 'WGS', 'Scaffolds', 'Genes', 'Proteins', 'Release Date',
      'Modify Date', 'Status'],
      dtype='object')
```

len(df_center.columns)

18

len(df_rejigged.columns)

... but luckily ...

df_center.index

```
Index(['JGI', 'The Arabidopsis Information Resource (TAIR)',  
      'Ocean University', 'US DOE Joint Genome Institute (JGI-PGF)', 'NCBI',  
      'Solanaceae Genomics Project',  
      'Leibniz Institute of Plant Genetics and Crop Plant Research', 'NCBI',  
      'MaizeGDB', 'NCBI',  
      ...  
      'University of Adelaide',  
      'Genome Sequencing Center (GSC) at Washington University (WashU) School of Medicine',  
      'Genome Sequencing Center (GSC) at Washington University (WashU) School of Medicine',  
      'Genome Sequencing Center (GSC) at Washington University (WashU) School of Medicine',  
      'Genome Sequencing Center (GSC) at Washington University (WashU) School of Medicine',  
      'Genome Sequencing Center (GSC) at Washington University (WashU) School of Medicine',  
      'Genome Sequencing Center (GSC) at Washington University (WashU) School of Medicine',  
      'Shanghai Jiao Tong University', 'Chung-Ang University',  
      'Sanger Institute'],  
      dtype='object', name='Center', length=37950)
```

We can also make our own index out of items in the data. This we do by assigning to the special variable `df.index` as we have just seen above, which lets us do things like....

Start again (just to be safe!): re-read data in with row number as default index

```
df = pd.read_csv('eukaryotes.txt', sep="\t", na_values=['-'])
```

We want to create a new index which is made up of the Organism name and the accession number

Check our command out first

```
df.apply(lambda x : "{} ({}).format(x['#Organism/Name'], x['BioSample  
Accession']), axis=1)
```

```
0      Emiliana huxleyi CCMP1516 (SAMN02744062)  
1      Arabidopsis thaliana (SAMN03081427)  
2      Neopyropia yezoensis (SAMN13316713)  
3      Glycine max (SAMN00002965)  
4      Medicago truncatula (SAMN08400029)  
...  
37945  Saccharomyces cerevisiae NC-02 (SAMN14365262)  
37946  Saccharomyces cerevisiae CLIB382 (SAMN00198998)  
37947  Saccharomyces cerevisiae (SAMN20308677)  
37948  Saccharomyces cerevisiae (SAMN26740367)  
37949  Saccharomyces cerevisiae (SAMEA3138293)  
Length: 37950, dtype: object
```

```
# OK, so let's make the index, as that seems to work
```

```
df.index = df.apply(lambda x : "{} ({}).format(x['#Organism/Name'],  
x['BioSample Accession']), axis=1)
```

```
df
```

```
                                     #Organism/Name ... Bi
Emiliana huxleyi CCMP1516 (SAMN02744062)      Emiliana huxleyi CCMP1516 ...
Arabidopsis thaliana (SAMN03081427)           Arabidopsis thaliana ...
Neopyropia yezoensis (SAMN13316713)          Neopyropia yezoensis ...
Glycine max (SAMN00002965)                    Glycine max ...
Medicago truncatula (SAMN08400029)           Medicago truncatula ...
...                                           ...
Saccharomyces cerevisiae NC-02 (SAMN14365262) Saccharomyces cerevisiae NC-02 ...
Saccharomyces cerevisiae CLIB382 (SAMN00198998) Saccharomyces cerevisiae CLIB382 ...
Saccharomyces cerevisiae (SAMN20308677)       Saccharomyces cerevisiae ...
Saccharomyces cerevisiae (SAMN26740367)       Saccharomyces cerevisiae ...
Saccharomyces cerevisiae (SAMEA3138293)       Saccharomyces cerevisiae ...

[37950 rows x 19 columns]
```

```
df.shape
```

```
(37950, 19)
```

```
df.index.value_counts()
```

```
Solanum verrucosum (SAMEA104077226)      19
Homo sapiens (SAMN03283347)               13
Taeniopygia guttata (SAMN09946140)        11
Mustela putorius (SAMEA5818864)           10
Homo sapiens (SAMN03255769)               9
..
Arthroderma sp. MBC 290 (SAMN30838493)    1
Lipomyces yarrowii (SAMN20341228)         1
[Candida] tunisiensis (SAMN20341277)      1
Beauveria australis (SAMN30838441)        1
Saccharomyces cerevisiae (SAMEA3138293)    1
Length: 35010, dtype: int64
```

```
df['TaxID'].head(3)
```

```
Emiliana huxleyi CCMP1516 (SAMN02744062)    280463
Arabidopsis thaliana (SAMN03081427)         3702
Neopyropia yezoensis (SAMN13316713)        2788
Name: TaxID, dtype: int64
```

```
df[0:3]['TaxID']
```

```
Emiliana huxleyi CCMP1516 (SAMN02744062)    280463
Arabidopsis thaliana (SAMN03081427)         3702
Neopyropia yezoensis (SAMN13316713)        2788
Name: TaxID, dtype: int64
```

```
df['TaxID'].values[0:3]
```

```
array([280463, 3702, 2788])
```

```
# Is it like a dictionary, got keys too!? No...
```

```
df['TaxID'].keys[0:3]
```

```
Traceback (most recent call last):  
  File "<stdin>", line 1, in <module>  
TypeError: 'method' object is not subscriptable
```

Selections using `iloc`

Single rows : first row

`df.iloc[0]`

```
#Organism/Name      Emiliana huxleyi CCMP1516
TaxID                280463
BioProject Accession PRJNA77753
BioProject ID        77753
Group                Protists
SubGroup              Other Protists
Size (Mb)             167.676
GC%                   64.5
Assembly Accession    GCA_000372725.1
Replicons             NaN
WGS                   AHAL01
Scaffolds             7795
Genes                 38549.0
Proteins              38554.0
Release Date          2013/04/19
Modify Date           2014/08/01
Status                Scaffold
Center                JGI
BioSample Accession    SAMN02744062
Name: Emiliana huxleyi CCMP1516 (SAMN02744062), dtype: object
```

Single rows : last row

`df.iloc[-1]`

```
#Organism/Name      Saccharomyces cerevisiae
TaxID                4932
BioProject Accession PRJNA13838
BioProject ID        13838
Group                Fungi
SubGroup              Ascomycetes
Size (Mb)             0.586761
GC%                   38.5921
Assembly Accession    GCA_000091065.2
Replicons              chromosome III:X59720.2; chromosome VI:D50617.1
WGS                   NaN
Scaffolds             2
Genes                 155.0
Proteins              298.0
Release Date          1992/03/16
Modify Date           2009/07/02
Status                Chromosome
Center                Sanger Institute
BioSample Accession    SAMEA3138293
Name: Saccharomyces cerevisiae (SAMEA3138293), dtype: object
```


Single columns : first column

df.iloc[:,0]

Emiliana huxleyi CCMP1516 (SAMN02744062)	Emiliana huxleyi CCMP1516
Arabidopsis thaliana (SAMN03081427)	Arabidopsis thaliana
Neopyropia yezoensis (SAMN13316713)	Neopyropia yezoensis
Glycine max (SAMN00002965)	Glycine max
Medicago truncatula (SAMN08400029)	Medicago truncatula
	...
Saccharomyces cerevisiae NC-02 (SAMN14365262)	Saccharomyces cerevisiae NC-02
Saccharomyces cerevisiae CLIB382 (SAMN00198998)	Saccharomyces cerevisiae CLIB382
Saccharomyces cerevisiae (SAMN20308677)	Saccharomyces cerevisiae
Saccharomyces cerevisiae (SAMN26740367)	Saccharomyces cerevisiae
Saccharomyces cerevisiae (SAMEA3138293)	Saccharomyces cerevisiae

Name: #Organism/Name, Length: 37950, dtype: object

Single columns : last column

df.iloc[:, -1]

Emiliana huxleyi CCMP1516 (SAMN02744062)	SAMN02744062
Arabidopsis thaliana (SAMN03081427)	SAMN03081427
Neopyropia yezoensis (SAMN13316713)	SAMN13316713
Glycine max (SAMN00002965)	SAMN00002965
Medicago truncatula (SAMN08400029)	SAMN08400029
	...
Saccharomyces cerevisiae NC-02 (SAMN14365262)	SAMN14365262
Saccharomyces cerevisiae CLIB382 (SAMN00198998)	SAMN00198998
Saccharomyces cerevisiae (SAMN20308677)	SAMN20308677
Saccharomyces cerevisiae (SAMN26740367)	SAMN26740367
Saccharomyces cerevisiae (SAMEA3138293)	SAMEA3138293

Name: BioSample Accession, Length: 37950, dtype: object

Multiple row and column selection examples

```
df.iloc[0:5]          # first five rows of dataframe
df.iloc[:, 0:2]       # first two columns of data frame with all rows
df.iloc[ [0,3,6,24], [0,5,6] ] # 1st, 4th, 7th, 25th row + 1st 6th 7th columns.
df.iloc[0:5, 5:8]     # first 5 rows and 5th 6th 7th columns
```

The actual outputs

df.columns

```
Index(['#Organism/Name', 'TaxID', 'BioProject Accession', 'BioProject ID',
      'Group', 'SubGroup', 'Size (Mb)', 'GC%', 'Assembly Accession',
      'Replicons', 'WGS', 'Scaffolds', 'Genes', 'Proteins', 'Release Date',
      'Modify Date', 'Status', 'Center', 'BioSample Accession'],
      dtype='object')
```

df.columns[5:8]

```
Index(['SubGroup', 'Size (Mb)', 'GC%'], dtype='object')
```

df.iloc[0:5, 5:8]

	SubGroup	Size (Mb)	GC%
Emiliana huxleyi CCMP1516 (SAMN02744062)	Other Protists	167.676	64.5000
Arabidopsis thaliana (SAMN03081427)	Land Plants	119.669	36.0529
Neopyropia yezoensis (SAMN13316713)	Other	107.591	64.8454

Give feedback
Ask for help

Glycine max (SAMN00002965)	Land Plants	978.942	35.1221
Medicago truncatula (SAMN08400029)	Land Plants	430.008	33.4462

📌 The Pandas `loc` indexer can be used with DataFrames for two different use cases:

- Selecting rows by label/index (as we have just done using `iloc`)
- Selecting rows with a boolean / conditional lookup (i.e. True or False)

The `loc` indexer is used with the same syntax as `iloc`:

```
data.loc[<row selection>, <column selection>]
```

Let's try it

```
df.loc['Glycine max (SAMN00002965)']
```

```
#Organism/Name          Glycine max
TaxID                   3847
BioProject Accession    PRJNA19861
BioProject ID           19861
Group                   Plants
SubGroup                Land Plants
Size (Mb)               978.942
GC%                     35.1221
Assembly Accession      GCA_000004515.5
Replicons                chromosome 1:NC_016088.4/CM000834.4; chromosom...
WGS                      ACUP04
Scaffolds                347
Genes                   59046.0
Proteins                 74248.0
Release Date            2010/01/05
Modify Date             2021/03/10
Status                  Chromosome
Center                  US DOE Joint Genome Institute (JGI-PGF)
BioSample Accession      SAMN00002965
Name: Glycine max (SAMN00002965), dtype: object
```

Let's try it with more than one

```
df.loc[ ['Glycine max (SAMN00002965)', 'Solanum lycopersicum (SAMN02981290)'] ]
```

```
                                #Organism/Name  ...  BioSample Accession
Glycine max (SAMN00002965)          Glycine max  ...          SAMN00002965
Solanum lycopersicum (SAMN02981290)  Solanum lycopersicum  ...          SAMN02981290

[2 rows x 19 columns]
```

Quick reminder, what columns do we have to choose from?

df.columns

```
Index(['#Organism/Name', 'TaxID', 'BioProject Accession', 'BioProject ID',
      'Group', 'SubGroup', 'Size (Mb)', 'GC%', 'Assembly Accession',
      'Replicons', 'WGS', 'Scaffolds', 'Genes', 'Proteins', 'Release Date',
      'Modify Date', 'Status', 'Center', 'BioSample Accession'],
      dtype='object')
```

We choose some columns

```
df.loc[ ['Glycine max (SAMN00002965)', 'Solanum lycopersicum
(SAMN02981290)'], ['TaxID', 'Size (Mb)', 'GC%', 'Genes'] ]
```

	TaxID	Size (Mb)	GC%	Genes
Glycine max (SAMN00002965)	3847	978.942	35.1221	59046.0
Solanum lycopersicum (SAMN02981290)	4081	827.963	35.6925	31219.0

Other way of doing it, same result!

```
df.loc[ ['Glycine max (SAMN00002965)', 'Solanum lycopersicum
(SAMN02981290)'], list(df.columns[ [1,6,7,12] ])]
```

	TaxID	Size (Mb)	GC%	Genes
Glycine max (SAMN00002965)	3847	978.942	35.1221	59046.0
Solanum lycopersicum (SAMN02981290)	4081	827.963	35.6925	31219.0

Other way of doing it, different column order!

```
df.loc[ ['Glycine max (SAMN00002965)', 'Solanum lycopersicum
(SAMN02981290)'], list(df.columns[ [6,7,1,12] ])]
```

	Size (Mb)	GC%	TaxID	Genes
Glycine max (SAMN00002965)	978.942	35.1221	3847	59046.0
Solanum lycopersicum (SAMN02981290)	827.963	35.6925	4081	31219.0

We choose a particular organism

```
df.loc[df['#Organism/Name']=='Venturia inaequalis']
```

	#Organism/Name	...	BioSample Accession
Venturia inaequalis (SAMN07816553)	Venturia inaequalis	...	SAMN07816553
Venturia inaequalis (SAMN07816628)	Venturia inaequalis	...	SAMN07816628
Venturia inaequalis (SAMN07816619)	Venturia inaequalis	...	SAMN07816619
Venturia inaequalis (SAMN03081268)	Venturia inaequalis	...	SAMN03081268
Venturia inaequalis (SAMN03080669)	Venturia inaequalis	...	SAMN03080669
...	...	...	...
Venturia inaequalis (SAMN07816556)	Venturia inaequalis	...	SAMN07816556
Venturia inaequalis (SAMN07816552)	Venturia inaequalis	...	SAMN07816552
Venturia inaequalis (SAMN07816568)	Venturia inaequalis	...	SAMN07816568
Venturia inaequalis (SAMN07816549)	Venturia inaequalis	...	SAMN07816549
Venturia inaequalis (SAMN07816547)	Venturia inaequalis	...	SAMN07816547

[88 rows x 19 columns]

Columns chosen, first three rows

```
df.loc[df['#Organism/Name']=='Venturia inaequalis', ['TaxID', 'Size
(Mb)', 'GC%', 'Genes'] ][0:3]
```

	TaxID	Size (Mb)	GC%	Genes
Venturia inaequalis (SAMN07816553)	5025	72.7916	43.1	NaN
Venturia inaequalis (SAMN07816628)	5025	51.7000	NaN	NaN
Venturia inaequalis (SAMN07816619)	5025	60.8229	NaN	NaN

# Columns chosen, first three rows, different way				
df.loc[df['#Organism/Name']=='Venturia inaequalis',['TaxID','Size (Mb)','GC%','Genes']].head(3)				
	TaxID	Size (Mb)	GC%	Genes
Venturia inaequalis (SAMN07816553)	5025	72.7916	43.1	NaN
Venturia inaequalis (SAMN07816628)	5025	51.7000	NaN	NaN
Venturia inaequalis (SAMN07816619)	5025	60.8229	NaN	NaN

# Looking for cows (<i>Bos taurus</i>)				
df.loc[df['#Organism/Name']=='Bos taurus']				
	#Organism/Name	TaxID	...	
Bos taurus (nan)	Bos taurus	9913	...	
Bos taurus (SAMN43074752)	Bos taurus	9913	...	Univers
Bos taurus (SAMN43074753)	Bos taurus	9913	...	Univers
Bos taurus (SAMEA115082612)	Bos taurus	9913	...	NORWEGIAN UNIVERSITY O
Bos taurus (SAMEA14214903)	Bos taurus	9913	...	IN
Bos taurus (SAMN32907642)	Bos taurus	9913	...	Seoul Nati
Bos taurus (SAMN16377989)	Bos taurus	9913	...	
Bos taurus (SAMN26241768)	Bos taurus	9913	...	Northwest
Bos taurus (SAMN37529654)	Bos taurus	9913	...	Yunnan Agricult
Bos taurus (SAMEA7047893)	Bos taurus	9913	...	THE R
Bos taurus (SAMEA7051353)	Bos taurus	9913	...	THE R
Bos taurus (SAMN16425342)	Bos taurus	9913	...	
Bos taurus (SAMN35152401)	Bos taurus	9913	...	Indian Institute of Science Educat
Bos taurus (SAMEA9551350)	Bos taurus	9913	...	the e
Bos taurus (SAMN02898106)	Bos taurus	9913	...	Cattle Genome Sequencing Internati
Bos taurus (SAMN02898106)	Bos taurus	9913	...	Center for Bioinformatics and Comp
Bos taurus (SAMN35642812)	Bos taurus	9913	...	
[17 rows x 19 columns]				

# Looking for a Polish cows				
df.loc[df['#Organism/Name']=='Polish cow']				
Empty DataFrame				
Columns: [#Organism/Name, TaxID, BioProject Accession, BioProject ID, Group, SubGroup, Siz				
Index: []				

None found thankfully!

📌 Some other quick cool things

```
df['Genes'].value_counts()
```

```
37.0      49
13.0      43
6907.0     5
6912.0     5
6911.0     5
..
15911.0    1
15454.0    1
15762.0    1
15393.0    1
155.0      1
Name: Genes, Length: 6150, dtype: int64
```

```
df[['#Organism/Name', 'Genes']].value_counts()
```

```
Aspergillus fumigatus      9481.0    5
Ophidiomyces ophidiicola  6892.0    5
Saccharomyces cerevisiae  6376.0    4
Aspergillus fumigatus      9570.0    4
                          9626.0    4
..
Eudypetes sclateri        13760.0    1
Eudypetes schlegeli       16825.0    1
Eudypetes robustus        16687.0    1
Eudypetes pachyrhynchus   18731.0    1
fungal sp. No.11243       9730.0    1
Length: 7170, dtype: int64
```

Aggregate

```
df['Genes'].agg('min')
```

```
1.0
```

```
df['Genes'].agg('max')
```

```
4736081.0
```

```
df['Genes'].agg(['max', 'min', 'mean', 'std'])
```

```
max      4.736081e+06
min      1.000000e+00
mean     1.822345e+04
std      8.392855e+04
Name: Genes, dtype: float64
```

```
type(df['Genes'].agg(['max', 'min', 'mean', 'std']))
```

<class 'pandas.core.series.Series'>

groupby

df.groupby('#Organism/Name').mean()

	TaxID	BioProject ID	Size (Mb)	GC%	Scaffolds	G
#Organism/Name						
Aaosphaeria arxii CBS 175.79	1450172.0	625762.0	38.901000	49.90	55.0	142
Aaosphaeria pasadenensis	3025424.0	800051.0	38.797400	49.70	120.0	
Abelmoschus esculentus	455045.0	985739.0	1194.600000	33.80	78.0	
Abeoforma whisleri	749232.0	360047.0	101.554000	31.40	71510.0	
Abisara bifasciata	2614024.0	606954.0	332.742000	32.30	754990.0	
...	...	...	...	...	...	
uncultured Thalassiosira	654899.0	955529.0	30.739600	51.60	5894.0	
uncultured Trebouxioephyceae	167235.0	809297.0	4.756060	73.00	2846.0	
uncultured Tremellales	227714.0	809291.0	14.909500	56.00	1783.0	
uncultured ciliate	187299.0	941656.0	17.869020	31.40	1146.2	
uncultured eukaryote	100272.0	850925.8	14.964724	59.32	908.6	

[17177 rows x 7 columns]

type(df.groupby('#Organism/Name').mean())

<class 'pandas.core.frame.DataFrame'>

df.groupby('#Organism/Name').mean().iloc[:,0:3]

	TaxID	BioProject ID	Size (Mb)
#Organism/Name			
Aaosphaeria arxii CBS 175.79	1450172.0	625762.0	38.901000
Aaosphaeria pasadenensis	3025424.0	800051.0	38.797400
Abelmoschus esculentus	455045.0	985739.0	1194.600000
Abeoforma whisleri	749232.0	360047.0	101.554000
Abisara bifasciata	2614024.0	606954.0	332.742000
...	...	...	...
uncultured Thalassiosira	654899.0	955529.0	30.739600
uncultured Trebouxioephyceae	167235.0	809297.0	4.756060
uncultured Tremellales	227714.0	809291.0	14.909500
uncultured ciliate	187299.0	941656.0	17.869020
uncultured eukaryote	100272.0	850925.8	14.964724

[17177 rows x 3 columns]

df.groupby('#Organism/Name').mean().iloc[:,0:3].head(10)

	TaxID	BioProject ID	Size (Mb)
#Organism/Name			
Aaosphaeria arxii CBS 175.79	1450172.0	625762.0	38.90100
Aaosphaeria pasadenensis	3025424.0	800051.0	38.79740
Abelmoschus esculentus	455045.0	985739.0	1194.60000
Abeoforma whisleri	749232.0	360047.0	101.55400
Abisara bifasciata	2614024.0	606954.0	332.74200
Abortiporus biennis	137743.0	444028.0	33.11860
Abramis brama	38527.0	946171.0	1018.32975
Abrostola tripartita	938171.0	715706.5	379.47800
Abrostola triplasia	254365.0	870861.5	361.63200
Abrus precatorius	3816.0	510631.0	347.23000

df.groupby('#Organism/Name').mean().iloc[0:14:,0:3]

	TaxID	BioProject ID	Size (Mb)
#Organism/Name			
Aaosphaeria arxii CBS 175.79	1450172.0	625762.0	38.90100
Aaosphaeria pasadenensis	3025424.0	800051.0	38.79740

Abelmoschus esculentus	455045.0	985739.0	1194.60000
Abeoforma whisleri	749232.0	360047.0	101.55400
Abisara bifasciata	2614024.0	606954.0	332.74200
Abortiporus biennis	137743.0	444028.0	33.11860
Abramis brama	38527.0	946171.0	1018.32975
Abrostola tripartita	938171.0	715706.5	379.47800
Abrostola triplasia	254365.0	870861.5	361.63200
Abrus precatorius	3816.0	510631.0	347.23000
Abrus pulchellus subsp. cantoniensis	858909.0	812623.5	381.26800
Abscondita cerata	2740424.0	699421.0	967.16200
Abscondita terminalis	2069292.0	556938.0	499.65300
Absidia glauca	4829.0	317991.0	48.74640


📌 Exercises: your turn!

Let's try and be tidy...

Make a directory within your homespace for today's work (if you didnt already!), and change directory to it.

```
mkdir ~/Exercises/Lecture17
```

```
cd ~/Exercises/Lecture17
```



Sending yet another email from the server

Today we are going to send yet another email from the bioinformatics server! You are going to use it to send me (and also to yourself) a short message about

your recent git usage.

Here is what you need to do, BEFORE you start any of today's exercises!

1. log in to the server
2. change directory to wherever you have your **Lecture Exercises git repository**
3. copy paste the following 3 lines of commands into your terminal window

```
DATE=$(date)
git log | grep "Date:" | egrep "Oct|Nov" > message_to_al_3
mail -s "${USER} ${HOSTNAME} ${DATE}" < message_to_al_3 aivens2@ed.ac.uk,${USER}@ed.ac.uk
```

Mail sent, now check your email!

Download the latest `eukaryotes.txt` equivalent file from NCBI (or alternatively, use the one provided in `/localdisk/data/BPSM/Lecture17`)

Use a pandas dataframe to address the following questions:

- 📡 how many fungal species have genomes bigger than 100Mb? What are their names?
 - 📡 how many of each Kingdom/group (plants, animals, fungi and protists) have been sequenced?
 - 📡 which *Heliconius* species genomes have been sequenced?
 - 📡 which sequencing centre has sequenced the most plant genomes? the most insect genomes?
 - 📡 add a column giving the number of proteins per gene. Which genomes have at least 10% more proteins than genes?
-
-

Possible answers

First import the libraries

```
import os, sys, re, numpy as np
import pandas as pd
```

Get the file

```
os.mkdir("${HOME}/Exercises/Lecture17")
```

```
os.chdir("${HOME}/Exercises/Lecture17")
```

If you opted for the new version, you could call it `eukaryotes.tsv` as its tab-separated values?

```
os.system("wget -qO eukaryotes.tsv
'ftp://ftp.ncbi.nlm.nih.gov/genomes/GENOME_REPORTS/eukaryotes.txt'")
```

Read the file into a dataframe

```
df = pd.read_csv('eukaryotes.tsv', sep="\t", na_values=['-'])
```

Create a new index which is made up of the species name and the accession number

```
df.index=df.apply(lambda x : "{} ({}).format(x['#Organism/Name'],
x['BioSample Accession']), axis=1)
```

What are the column headers

```
df.columns
```

```
Index(['#Organism/Name', 'TaxID', 'BioProject Accession', 'BioProject ID',
      'Group', 'SubGroup', 'Size (Mb)', 'GC%', 'Assembly Accession',
      'Replicons', 'WGS', 'Scaffolds', 'Genes', 'Proteins', 'Release Date',
      'Modify Date', 'Status', 'Center', 'BioSample Accession'],
      dtype='object')
```

How many fungal species have genomes bigger than 100Mb?

use two conditions joined with `&`

```
len( df[ (df['Group'] == 'Fungi') & (df['Size (Mb)'] > 100) ] )
```

218

or we could use `apply`

```
len(df[df.apply(lambda x : x['Group'] == 'Fungi' and x['Size (Mb)'] > 100,
axis=1)])
```

218

What are their names?

Start with the same query as before, then get the names column

```
big_fungi = df[(df['Group'] == 'Fungi') & (df['Size (Mb)'] > 100)]
```

Transform the series object into a simple list of strings

```
list(big_fungi['#Organism/Name'])
```

```
['Blumeria graminis f. sp. triticales', 'Puccinia triticina', 'Tuber melanosporum Mel28',  
...  
'Rhizophagus irregularis DAOM 197198w', 'Rhizophagus irregularis DAOM 197198w', 'Rhizophagus irregularis']
```

Nicer when it is sorted!

```
sorted(list(big_fungi['#Organism/Name']))
```

```
['Acaulospora colombiana', 'Acaulospora morrowiae', 'Albatrellus ellisii',  
...  
'Verticillium longisporum', 'Wilcoxina mikolae CBS 423.85', 'Zopfia rhizophila CBS 207.26']
```

How many genomes of each major group (plants, animals, fungi and protists) have been sequenced, and how many unique organisms?

We can figure out the answer for a single group

```
len(df[df['Group'] == 'Plants'])
```

4390

```
len(df[df['Group'] == 'plants'])
```

0

```
len(df[df['Group'].str.lower() == 'plants'])
```

4390

```
len(df[df['Group'].str.lower() == 'PLANTS'])
```

0

```
len(df[df['Group'].str.upper() == 'PLANTS'])
```

4390

Remember that we are in Python3, so we can use a normal loop

```
for Group in ['Plants', 'Animals', 'Fungi', 'Protists']:  
    count = len(df[df['Group'] == Group])
```

To count unique ones, we could get the names and turn them into a set

```
count_unique = len(set(df[df['Group'] == Group]['#Organism/Name']))
```

OR we could use the `drop_duplicates` method which may be clearer to read

```
count_unique = len(df[df['Group'] == Group].drop_duplicates('#Organism/Name'))  
print("{} genomes for {} ({} unique)".format(count, Group, count_unique))
```

```
4390 genomes for Plants (1912 unique)
14495 genomes for Animals (8190 unique)
16876 genomes for Fungi (6060 unique)
1987 genomes for Protists (871 unique)
```

Which *Heliconius* species genomes have been sequenced, and how many scaffolds is each assembly in?

here we HAVE to use apply

```
hel = df[df.apply(lambda x : x['#Organism/Name'].startswith('Heliconius'), axis=1)]
hel[ ['#Organism/Name', 'Scaffolds'] ]
```

	#Organism/Name	Scaffolds
Heliconius cydno galanthus (SAMN25599327)	Heliconius cydno galanthus	1341
Heliconius melpomene melpomene (SAMEA2272046)	Heliconius melpomene melpomene	4309
Heliconius numata (SAMN16746774)	Heliconius numata	15243
Heliconius pachinus (SAMEA3670564)	Heliconius pachinus	283902
Heliconius ismenius (SAMEA3670537)	Heliconius ismenius	160084
...	...	...
Heliconius melpomene amaryllis (SAMEA1094479)	Heliconius melpomene amaryllis	104829
Heliconius melpomene amaryllis (SAMEA1094482)	Heliconius melpomene amaryllis	46387
Heliconius melpomene amaryllis (SAMEA1094488)	Heliconius melpomene amaryllis	116151
Heliconius melpomene aglaope (SAMEA1094472)	Heliconius melpomene aglaope	49473
Heliconius melpomene aglaope (SAMEA1094478)	Heliconius melpomene aglaope	105664

[182 rows x 2 columns]

Which center has sequenced the most plant genomes?

```
df[df['Group'] == 'Plants']['Center'].value_counts().head()
```

USDA-ARS	314
Leibniz Institute of Plant Genetics and Crop Plant Research	185
WELLCOME SANGER INSTITUTE	167
Iridian Genomes	151
NCBI	141

Name: Center, dtype: int64

Which center has sequenced the most insect genomes?

```
(df['Group'] == 'Insects').value_counts()
```

```
False    26070
Name: Group, dtype: int64
```

Oops!

```
(df['SubGroup'] == 'Insects').value_counts()
```

```
False    32858
True      5092
Name: SubGroup, dtype: int64
```

```
df[df['SubGroup'] == 'Insects']['Center'].value_counts().head()
```

WELLCOME SANGER INSTITUTE	1637
Florida Museum of Natural History	816
Stanford University	274
NCBI	249
University of Cambridge, UK	141
Name: Center, dtype: int64	

Add a column giving the number of proteins per gene

first, we do the calculation

df['Proteins'] / df['Genes']

Emiliana huxleyi CCMP1516 (SAMN02744062)	1.000130
Arabidopsis thaliana (SAMN03081427)	1.259821
Neopyropia yezoensis (SAMN13316713)	NaN
Glycine max (SAMN00002965)	1.257460
Medicago truncatula (SAMN08400029)	1.078153
...	
Saccharomyces cerevisiae NC-02 (SAMN14365262)	NaN
Saccharomyces cerevisiae CLIB382 (SAMN00198998)	NaN
Saccharomyces cerevisiae (SAMN20308677)	NaN
Saccharomyces cerevisiae (SAMN26740367)	NaN
Saccharomyces cerevisiae (SAMEA3138293)	1.922581
Length: 37950, dtype: float64	

We can assign this to a new column

df['Proteins per gene'] = df['Proteins'] / df['Genes']

Show the results

df[['#Organism/Name', 'Group', 'Proteins per gene']].head()

	#Organism/Name	Group	Proteins
Emiliana huxleyi CCMP1516 (SAMN02744062)	Emiliana huxleyi CCMP1516	Protists	
Arabidopsis thaliana (SAMN03081427)	Arabidopsis thaliana	Plants	
Neopyropia yezoensis (SAMN13316713)	Neopyropia yezoensis	Other	
Glycine max (SAMN00002965)	Glycine max	Plants	
Medicago truncatula (SAMN08400029)	Medicago truncatula	Plants	

Which genomes have at least 10% more proteins than genes?

df[df['Proteins per gene'] >= 1.1][['#Organism/Name', 'Genes', 'Proteins', 'Proteins per gene']].head()

	#Organism/Name	Genes	Proteins	Proteins per
Arabidopsis thaliana (SAMN03081427)	Arabidopsis thaliana	38311.0	48265.0	1.2
Glycine max (SAMN00002965)	Glycine max	59046.0	74248.0	1.2
Solanum lycopersicum (SAMN02981290)	Solanum lycopersicum	31219.0	37604.0	1.2
Zea mays (SAMEA5569141)	Zea mays	49897.0	57578.0	1.1
Aedes aegypti (SAMN07177802)	Aedes aegypti	19623.0	28317.0	1.4

df[df['Proteins per gene'] >= 1.1][['#Organism/Name', 'Genes', 'Proteins', 'Proteins per gene']]

Arabidopsis thaliana (SAMN03081427)	Arabidops
Glycine max (SAMN00002965)	
Solanum lycopersicum (SAMN02981290)	Solanum 1
Zea mays (SAMEA5569141)	

#Or

```
Aedes aegypti (SAMN07177802)
...
Fusarium oxysporum f. sp. vasinfectum 25433 (SA...
Fusarium oxysporum f. sp. raphani 54005 (SAMN02...
Fusarium oxysporum f. sp. conglutinans race 2 5...
Fusarium oxysporum NRRL 32931 (SAMN02981331)
Saccharomyces cerevisiae (SAMEA3138293)

Fusarium oxysporum f. sp. vasinfectum 25433 (SA...
Fusarium oxysporum f. sp. raphani 54005 (SAMN02...
Fusarium oxysporum f. sp. conglutinans race 2 5...
Fusarium oxysporum NRRL 32931 (SAMN02981331)
Saccharomyces cerevisiae (SAMEA3138293)

[1070 rows x 4 columns]
```

```
len(df[ df['Proteins per gene'] >= 1.1])
```

```
1070
```

```
# More than 2.5 proteins per gene...
```

```
df[df['Proteins per gene'] >= 2.5][ ['#Organism/Name',
'Genes','Proteins','Proteins per gene'] ]
```

	#Organism/Name	Genes
Gallus gallus (SAMN15960293)	Gallus gallus	25638.0
Pan troglodytes (SAMN30216104)	Pan troglodytes	41815.0
Citrus sinensis (SAMN19611724)	Citrus sinensis	55825.0
Triticum turgidum subsp. durum (SAMEA104312462)	Triticum turgidum subsp. durum	63993.0
Enteromyxum leei (SAMN03701405)	Enteromyxum leei	5.0
Linum tenue (SAMEA110422887)	Linum tenue	52052.0
Phyllostomus discolor (SAMN14734468)	Phyllostomus discolor	21058.0
Citrus sinensis (SAMN19611724)	Citrus sinensis	55880.0
Schistosoma japonicum (SAMN19770425)	Schistosoma japonicum	9768.0
Engystomops pustulosus (SAMN13066191)	Engystomops pustulosus	18426.0
Ascaphus truei (SAMN12123227)	Ascaphus truei	17155.0
Gallus gallus (SAMN15960293)	Gallus gallus	25612.0
Zea mays (SAMN04296295)	Zea mays	39320.0

📈git and GitHub time!

```
# Local repository actions (the absolute minimum set of commands required here...)
```

```
git st
```

```
On branch master
Changes not staged for commit:
  (use "git add/rm <file>..." to update what will be committed)
  (use "git restore <file>..." to discard changes in working directory)
    deleted:    ../Lecture04/CodeFiles.tar.gz
    modified:   ../Lecture06/run_this_script_v1

Untracked files:
  (use "git add <file>..." to include in what will be committed)
    ../BN.out
    ../BN.sh
    ../Lecture04/Als_ICA/
    ../Lecture04/all_the_things_I_did
    ../Lecture04/outs/
    ../Lecture05/Country.Mozambique.details
    ../Lecture07/Cosmoscarta.nuc.fa
    ../Lecture07/Cosmoscarta.nuc.gis
```

```
../Lecture07/myfileofaccessions.txt
../Lecture12/AJ223353.fasta
../Lecture12/AJ223353.genbank
../Lecture12/AJ223353_noheader.fasta
../Lecture12/AJ223353_noheader.fasta2
../Lecture12/AJ223353_noheader.fasta3
../Lecture12/All_exons.fasta
../Lecture12/All_noncodings.fasta
../Lecture12/local_exon01.fasta
../Lecture12/local_exon02.fasta
../Lecture12/local_intron01.fasta
../Lecture12/my_dna.txt
../Lecture12/plain_genomic_seq.txt
../Lecture12/quick_blast_check.out
../Lecture12/remote_exon01.fasta
../Lecture12/remote_noncoding01.fasta
../Lecture12/remote_noncoding02.fasta
../Lecture12/shutilversion.txt
../Lecture13/100_199/
../Lecture13/200_299/
../Lecture13/300_399/
../Lecture13/400_499/
../Lecture13/500_599/
../Lecture13/600_699/
../Lecture13/700_799/
../Lecture13/800_899/
../Lecture13/900_999/
../Lecture13/Cleaned_sequences.txt
../Lecture13/Cleaned_sequences2.txt
../Lecture13/Cleaned_sequences3.txt
../Lecture13/Exercise2_coding_sequence.fasta
../Lecture13/Lecture13.tar.gz
../Lecture13/dna_files/
../Lecture13/exons.txt
../Lecture13/genomic_dna2.txt
../Lecture13/input.txt
../Lecture13/remote_exon01.fasta
../Lecture14/__pycache__/
../Lecture14/data.csv
../Lecture16/__pycache__/
./
```

no changes added to commit (use "git add" and/or "git commit -a")

ls -l

```
All_the_Arabidopsis_ones.tsv
eukaryotes.txt
myPandacommands.py
playtime
```

You might want to add other files?

git add *.py

warning: in the working copy of 'Lecture17/myPandacommands.py', CRLF will be replaced by LF the next time Git touches it

git commit -m "Script from the python pandas lecture"

```
[master 097800e] Script from the python pandas lecture
1 file changed, 96 insertions(+)
create mode 100644 Lecture17/myPandacommands.py
```

Push our local repository to GitHub

git remote

AllLectures
Lecture07
origin

git push AllLectures master

```
Enumerating objects: 5, done.
Counting objects: 100% (5/5), done.
Delta compression using up to 256 threads
Compressing objects: 100% (3/3), done.
Writing objects: 100% (4/4), 1.67 KiB | 857.00 KiB/s, done.
Total 4 (delta 1), reused 0 (delta 0), pack-reused 0
remote: Resolving deltas: 100% (1/1), completed with 1 local object.
To https://github.com/B123456-2121/BPSM_LectureStuff.git
 1d2319e..097800e  master -> master
```

git log

```
097800e (HEAD -> master, AllLectures/master) Script from the python pandas lecture
1d2319e Scripts from the python dictionaries lecture
c49df45 Scripts from the python functions lecture
3b8bf41 My scripts from the python conditionals exercises Py04
8f58f4c Scripts from the python lists and loops lecture 3
e444ef8 Scripts from the second BPSM python lecture
6d6f5e7 My python scripts from the first BPSM python lecture
eae65c6 (Lecture07/master) edirect seqs related stuff added
eea4f5e (Lec_6_branch) Kept bash script output files from Lecture 06
fba5cc0 Added run_this_script_v1 from BPSM Lecture06
ce2c15a nematode BLASTDB stuff added
19cdd2d Added the example_people_data.tsv file from Lecture03
137c3db Decided to keep blastoutput2.out
08d379b Decided to keep myinputfile.tsv
c9b5b18 First attempt at parsing BLAST data
bbf825c My first awk script, Lecture05
68f27fb There were conflicts with someotherfile.sh, kept all changes
83629cb (Version1.2) Final version of the someotherfile.sh shell script prior to merge
0b5425b (tag: v2.0) Updated contents of tar, coolscript
5d8aeea I am not sure cool script is actually any good
bbd36e7 Added my_cool_script
ea60357 Third version of the someotherfile.sh shell script
b6ae4a7 Updated version of tar
ce230e1 (tag: v1.2) Second version of the someotherfile.sh shell script
a5add7c Added the someotherfile.sh shell script
101b700 Additional motif files, and updated tar file
d21cb3d Adding intial tar file
59d1756 CodeFiles all done and tar zipped
6dd85a8 First file added
```

Have your say...

The "Postbox of Truth": is your chance to give anonymous "in course" feedback/comments. Your comments are completely anonymous and confidential, so say what you think, irrespective of whether the comments are positive or otherwise! You can use this multiple times too: I use it throughout the course to enable you to give context-relevant feedback/make suggestions.

Just click the image!



Sources used